

CO

Classification using MNIST.ipynb

☆

文件 修改 视图 插入 代码执行程序 工具 帮助 已保存所有更改

评论 分享 设置 S

+ 代码 + 文本

RAM 磁盘

修改

Chapter 3 – Classification

This notebook contains all the sample code and solutions to the exercises in chapter 3.

Open in Colab Open in Kaggle

Setup

First, let's import a few common modules, ensure Matplotlib plots figures inline and prepare a function to save the figures. We also check that Python 3.5 or later is installed (although Python 2.x may work, it is deprecated so we strongly recommend you use Python 3 instead), as well as Scikit-Learn ≥0.20.

```
[1] # Python ≥3.5 is required
import sys
assert sys.version_info >= (3, 5)

# Is this notebook running on Colab or Kaggle?
IS_COLAB = "google.colab" in sys.modules
IS_KAGGLE = "kaggle_secrets" in sys.modules

# Scikit-Learn ≥0.20 is required
```

```
# Common imports
import numpy as np
import os

# to make this notebook's output stable across runs
np.random.seed(42)

# To plot pretty figures
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rcParams['axes', labelsizes=14]
mpl.rcParams['xtick', labelsizes=12]
mpl.rcParams['ytick', labelsizes=12]

# Where to save the figures
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "classification"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

MNIST

Warning: since Scikit-Learn 0.24, `fetch_openml()` returns a Pandas `DataFrame` by default. To avoid this and keep the same code as in the book, we use `as_frame=False`.

```
[2] from sklearn.datasets import fetch_openml
mnist = fetch_openml("mnist_784", version=1, as_frame=False)
mnist.keys()

dict_keys(['data', 'target', 'frame', 'categories', 'feature_names', 'target_names', 'DESCR', 'details', 'url'])
```

```
[3] X, y = mnist["data"], mnist["target"]
X.shape

(70000, 784)
```

```
[4] y.shape

(70000,)
```

```
[5] 28 * 28


784
```

```
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt

some_digit = X[0]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap=mpl.cm.binary)
plt.axis("off")

save_fig("some_digit_plot")
plt.show()
```

Saving figure some_digit_plot



```
'5'
```

```
[8] y = y.astype(np.uint8)
```

```
[9] def plot_digit(data):
    image = data.reshape(28, 28)
    plt.imshow(image, cmap = mpl.cm.binary,
               interpolation="nearest")
    plt.axis("off")
```

```
# EXTRA
def plot_digits(instances, images_per_row=10, **options):
    size = 28
    images_per_row = min(len(instances), images_per_row)
    # This is equivalent to n_rows = ceil(len(instances) / images_per_row):
    n_rows = (len(instances) - 1) // images_per_row + 1


    # Append empty images to fill the end of the grid, if needed:
    n_empty = n_rows * images_per_row - len(instances)
    padded_instances = np.concatenate([instances, np.zeros((n_empty, size * size))], axis=0)

    # Reshape the array so it's organized as a grid containing 28x28 images:
    image_grid = padded_instances.reshape((n_rows, images_per_row, size, size))

    # Combine axes 0 and 2 (vertical image grid axis, and vertical image axis),
```

```
example_images = X[:100]
plot_digits(example_images, images_per_row=10)
save_fig("more_digits_plot")
plt.show()
```

Saving figure more_digits_plot



5026785904
6746807831

[12] y[0]

5

[13] X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]

Training a Binary Classifier

[14] y_train_5 = (y_train == 5)
y_test_5 = (y_test == 5)

Note: some hyperparameters will have a different default value in future versions of Scikit-Learn, such as max_iter and tol. To be future-proof, we explicitly set these hyperparameters to their future default values. For simplicity, this is not shown in the book.

[15] from sklearn.linear_model import SGDClassifier

sgd_clf.fit(X_train, y_train_5)

SGDClassifier(random_state=42)

[16] sgd_clf.predict([some_digit])

array([True])

[17] from sklearn.model_selection import cross_val_score
cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")

array([0.95035, 0.96035, 0.9604])

Performance Measures

Measuring Accuracy Using Cross-Validation

[18] from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone

skfolds = StratifiedKFold(n_splits=3, shuffle=True, random_state=42)

for train_index, test_index in skfolds.split(X_train, y_train_5):

X_train_folds = X_train[train_index]
y_train_folds = y_train_5[train_index]
X_test_fold = X_train[test_index]
y_test_fold = y_train_5[test_index]

clone_clf.fit(X_train_folds, y_train_folds)
y_pred = clone_clf.predict(X_test_fold)
n_correct = sum(y_pred == y_test_fold)
print(n_correct / len(y_pred))

0.9669
0.91625
0.96785

Note: shuffle=True was omitted by mistake in previous releases of the book.

from sklearn.base import BaseEstimator
class Never5Classifier(BaseEstimator):
def fit(self, X, y=None):
pass
def predict(self, X):
return np.zeros((len(X), 1), dtype=bool)

never_5_clf = Never5Classifier()
cross_val_score(never_5_clf, X_train, y_train_5, cv=3, scoring="accuracy")

array([0.91125, 0.90855, 0.90915])

+ 代码

+ 文本

✓ 磁盘

修改

^

Confusion Matrix

[21] from sklearn.model_selection import cross_val_predict
y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)
[22] from sklearn.metrics import confusion_matrix
confusion_matrix(y_train_5, y_train_pred)
array([[53892, 687],
[1891, 3530]])
y_train_perfect_predictions = y_train_5 # pretend we reached perfection
confusion_matrix(y_train_5, y_train_perfect_predictions)
array([[54579, 0],
[0, 5421]])

Precision and Recall

[24] from sklearn.metrics import precision_score, recall_score

+ 代码

+ 文本

✓ 磁盘

修改

^

0.8370879772350012
[25] cm = confusion_matrix(y_train_5, y_train_pred)
cm[1, 1] / (cm[0, 1] + cm[1, 1])
0.8370879772350012
[26] recall_score(y_train_5, y_train_pred)
0.6511713705958311
[27] cm[1, 1] / (cm[1, 0] + cm[1, 1])
0.6511713705958311
[28] from sklearn.metrics import f1_score
f1_score(y_train_5, y_train_pred)
0.7325171197343846
[29] cm[1, 1] / (cm[1, 1] + (cm[1, 0] + cm[0, 1]) / 2)
0.7325171197343847

Precision/Recall Trade-off


[30] y_scores = sgd_clf.decision_function([some_digit])
y_scores
array([2164.22030239])
[31] threshold = 0
y_some_digit_pred = (y_scores > threshold)
[32] y_some_digit_pred
array([True])
[33] threshold = 8000
y_some_digit_pred = (y_scores > threshold)
y_some_digit_pred
array([False])
[34] y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,
method="decision_function")
[35] from sklearn.metrics import precision_recall_curve

```
[36] def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision", linewidth=2)
    plt.plot(thresholds, recalls[:-1], "g-", label="Recall", linewidth=2)
    plt.legend(loc="center right", fontsize=16) # Not shown in the book
    plt.xlabel("Threshold", fontsize=16) # Not shown
    plt.grid(True) # Not shown
    plt.axis([-50000, 50000, 0, 1]) # Not shown

    recall_90_precision = recalls[np.argmax(precisions >= 0.90)]
    threshold_90_precision = thresholds[np.argmax(precisions >= 0.90)]

    plt.figure(figsize=(8, 4))
    plot_precision_recall_vs_threshold(precisions, recalls, thresholds) # No
    plt.plot([threshold_90_precision, threshold_90_precision], [0, 0.9], "r:") # Not shown
    plt.plot([-50000, threshold_90_precision], [0.9, 0.9], "r:") # Not shown
    plt.plot([-50000, threshold_90_precision], [recall_90_precision, recall_90_precision], "r:") # Not shown
    plt.plot([threshold_90_precision, [0.9], "ro") # Not shown
    plt.plot([threshold_90_precision, [recall_90_precision], "ro") # Not shown
    save_fig("precision_recall_vs_threshold_plot") # Not shown
    plt.show()
```

Saving figure precision_recall_vs_threshold_plot




```
[37] (y_train_pred == (y_scores > 0)).all()

True
```

```
def plot_precision_vs_recall(precisions, recalls):
    plt.plot(recalls, precisions, "b-", linewidth=2)
    plt.xlabel("Recall", fontsize=16)
    plt.ylabel("Precision", fontsize=16)
    plt.axis([0, 1, 0, 1])
    plt.grid(True)
```

```
plt.plot([recall_90_precision, recall_90_precision], [0, 0.9], "r:")
plt.plot([0, recall_90_precision], [0.9, 0.9], "r:")
plt.plot([recall_90_precision, [0.9], "ro")
save_fig("precision_vs_recall_plot")
plt.show()
```

Saving figure precision_vs_recall_plot



```
[39] threshold_90_precision = thresholds[np.argmax(precisions >= 0.90)]

[40] threshold_90_precision

3370.0194991439557

[41] y_train_pred_90 = (y_scores >= threshold_90_precision)

[42] precision_score(y_train_5, y_train_pred_90)

0.9000345901072293

[43] recall_score(y_train_5, y_train_pred_90)

0.4799852425751706
```

▼ The ROC Curve



```
[44] from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)
```

```
def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--') # dashed diagonal
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate (Fall-Out)', fontsize=16) # Not shown
    plt.ylabel('True Positive Rate (Recall)', fontsize=16) # Not shown
    plt.grid(True)

plt.figure(figsize=(8, 6))
plot_roc_curve(fpr, tpr)
fpr_90 = fpr[np.argmax(tpr >= recall_90_precision)] # Not shown
plt.plot([fpr_90, fpr_90], [0., recall_90_precision], 'r:') # Not shown
plt.plot([0.0, fpr_90], [recall_90_precision, recall_90_precision], 'r:') # Not shown
plt.plot([fpr_90], [recall_90_precision], 'ro') # Not shown
save_fig('roc_curve_plot') # Not shown
plt.show()
```

Saving figure roc_curve_plot



```
from sklearn.metrics import roc_auc_score

roc_auc_score(y_train_5, y_scores)

0.9604938554008616
```

Note: we set `n_estimators=100` to be future-proof since this will be the default value in Scikit-Learn 0.22.

```
[47] from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(n_estimators=100, random_state=42)
y_proba_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3,
```

0秒

[48] y_scores_forest = y_probas_forest[:, 1] # score = proba of positive class
fpr_forest, tpr_forest, thresholds_forest = roc_curve(y_train_5, y_scores_forest)

0秒

recall_for_forest = tpr_forest[np.argmax(fpr_forest >= fpr_90)]

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, "b:", linewidth=2, label="SGD")
plot_roc_curve(fpr_forest, tpr_forest, "Random Forest")
plt.plot([fpr_90, fpr_90], [0., recall_90_precision], "r:")
plt.plot([0.0, fpr_90], [recall_90_precision, recall_90_precision], "r:")
plt.plot([fpr_90, [recall_90_precision], "ro")
plt.plot([fpr_90, fpr_90], [0., recall_for_forest], "r:")
plt.plot([fpr_90, [recall_for_forest], "ro")
plt.grid(True)
plt.legend(loc="lower right", fontsize=16)
save_fig("roc_curve_comparison_plot")
plt.show()

Saving figure roc_curve_comparison_plot

0秒

roc_auc_score(y_train_5, y_scores_forest)

0.9983436731328145

1分钟

[51] y_train_pred_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3)
precision_score(y_train_5, y_train_pred_forest)

0.9905083315756169

0秒

[52] recall_score(y_train_5, y_train_pred_forest)

0.8662608374838591

Multiclass Classification

0秒

[53] from sklearn.svm import SVC

svm_clf = SVC(gamma="auto", random_state=42)
svm_clf.fit(X_train[:1000], y_train[:1000]) # y_train, not y_train_5
svm_clf.predict([some_digit])

array([5], dtype=uint8)

0秒

[54] some_digit_scores = svm_clf.decision_function([some_digit])
some_digit_scores

array([[2.81585438, 7.09167958, 3.82972099, 0.79365551, 5.8885703 ,
 9.29718395, 1.79862509, 8.10392157, -0.228207 , 4.83753243]])

0秒

[55] np.argmax(some_digit_scores)

5

0秒

[56] svm_clf.classes_

array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9], dtype=uint8)

0秒

[57] svm_clf.classes_[5]

5